



AIFMRM

AFRICAN INSTITUTE OF FINANCIAL MARKETS AND RISK MANAGEMENT



Deep Hedging under Black–Scholes and Heston Dynamics

by

Micaela Marais, Glasson Osborne, Kaiden Wessels & Siphелеle Nkosi

Supervised by Prof. Peter Ouwehand

Contents


1. Minimal Task
2. Parameter Robustness Study
3. Heston Extension

The hedge is learned through terminal portfolio error

$$\text{HE} = \pi + \underbrace{\sum_{n=0}^{N-1} \delta_n (S_{t_{n+1}} - S_{t_n})}_{\text{self-financing gains}} - (S_T - K)^+,$$
$$\min_{\pi, \theta} \mathbb{E}[\text{HE}^2].$$


The analytic delta is never used as a training target.

- ▶ $N = 125$ discrete hedge dates in the baseline.
- ▶ $\delta_n = f_\theta(x_n)$ is adapted to information at t_n .
- ▶ Negative HE is a seller shortfall; reported seller loss is $L = -\text{HE}$.
- ▶ With $r = 0$, the submitted experiments use undiscounted gains.



Source: Submitted report, Secs. 1.1, 2.2.2 and 3.1; Eqns. (1.1)–(1.2).

The implementation is differentiable end to end




Supervision is financial: payoff minus self-financing wealth.

Boundary: Function names and path construction are verified from report Appendix C; the separate executable notebook bundle is absent from this repository.

Source: Submitted report, Algorithm A.6.1 and Listings C.1–C.3. Shape contract follows the displayed array construction and shared network input.

Call economics guide architecture selection



one shared map at all 125 dates

- ▶ **Log-moneyness** is scale-normalised and financially natural.
- ▶ **Weight sharing** exploits Markov and time-homogeneous structure.
- ▶ **Sigmoid** enforces the call-delta range $0 < \delta < 1$.
- ▶ A scalar premium is learned jointly with the hedge.

Source: Submitted report, Secs. 3.1.2–3.1.3; Tables 3.1 and A.3–A.6.

Candidate	Test RMSE	Time (s)
Shared norm. 64×3 , tanh	0.007891	103.0
Shared norm. 64×2 , ReLU	0.007912	93.2
Time-separated 16×1	0.008083	786.9

Selected on mean validation loss across three seeds; close alternatives remained.

Boundary: The model was selected, not uniquely dominant: the ReLU candidate won one seed.

The fair benchmark matches the grid and objective

No hedge

raw option risk

Black–Scholes delta

continuous-time rule sampled on the grid

Discrete-time MSE-optimal

conditional projection on the same grid

Degree-2 polynomial

deliberately low-capacity approximation

Neural hedge

terminal-MSE training without delta labels

same held-out paths

$$\delta_n^{\text{DT}} = \frac{\text{Cov}(g(S_T), \Delta S_n \mid S_{t_n})}{\text{Var}(\Delta S_n \mid S_{t_n})}$$

The comparison asks whether the network recovers the finite-grid solution—not whether it beats correct Black–Scholes theory.

Source: Submitted report, Sec. 3.1.1 and Appendix A.4; Schweizer (1995, 2001).

The neural hedge recovers the finite-grid benchmark

Strategy	RMSE
No hedge	0.227 257
Black–Scholes delta	0.007 810
Discrete-time MSE-optimal	0.007 810
Degree-2 polynomial	0.022 183
Neural hedge	0.007 841

Boundary: These are final-benchmark values from an independent retraining, not architecture-search rows.

0.4%

neural RMSE above the discrete-time
benchmark


0.163972

learned premium
vs. analytic 0.164111

**Near equality is the validation
result we wanted.**

Source: Submitted report, Table 3.2 and Sec. 3.1.4.

The learned hedge has the right financial shape



- ▶ monotone in moneyness;
- ▶ close to both analytic curves;
- ▶ nearly indistinguishable on representative OTM, ATM and ITM paths.

The network learned the expected hedge function, not only a low aggregate loss.

Source: Economic reading of submitted report Figures 3.3 and 3.5; schematic above is not a data redraw.

Parameter Robustness under Retraining


- ▶ The same architecture is retrained from scratch at each setting.
- ▶ Parameters are varied one at a time around the baseline $(K, \sigma, T, N) = (0.9, 0.4, 0.5, 125)$.

$$K \in \{0.8, 0.9, 1.0, 1.1, 1.2\}$$

$$\sigma \in \{0.2, 0.4, 0.6\}$$

$$T \in \{0.25, 0.5, 1.0\}$$

$$N \in \{30, 60, 125, 250\}$$



Mean neural-to-optimal RMSE

ratio


1.049

(Range 1.009–1.099)

- ▶ Neural hedge remains at most 10% above the discrete-time optimum
- ▶ Largest relative gap occurs at $N = 250$

Generalisation Failure of the Single-Scenario Hedger

- ▶ One model is trained at $(K^*, \sigma^*) = (0.9, 0.4)$ and evaluated elsewhere **without retraining**.



Horizontal axis: (K, σ) combinations.

Why does it fail?

1. **Single training scenario:** No variation from which to learn dependence on K or σ .
2. **Strike is hardcoded:** K^* is built into log-moneyness.
3. **Volatility is missing:** The network has no input through which it can respond to a change in σ .

The Parameter-Conditioned Hedger

	Original network	Conditioned network
Inputs	$\left(\log \frac{S}{K}, \frac{\tau}{T}\right)$	$\left(\log \frac{S}{K}, \frac{\tau}{T}, K, \sigma\right)$
Training	Fixed $K = 0.9$, $\sigma = 0.4$	$K \sim U(0.7, 1.1)$, $\sigma \sim U(0.2, 0.6)$
Paths	100 000	200 000
Premium	Learned jointly	Analytic Black–Scholes premium

Interpolation

Inside the training range

$$K \in [0.7, 1.1], \quad \sigma \in [0.2, 0.6].$$

- ▶ Tested across nine parameter combinations.
- ▶ Maximum RMSE gap: ≤ 0.00018

Extrapolation

Outside the training range

$$K \in \{0.6, 1.3\}, \quad \sigma \in \{0.1, 0.8\}.$$

- ▶ Strike extrapolation remains within 0.00004 of the optimum.
- ▶ Worst vol. gap at $\sigma = 0.1$: $\approx 3 \times$ the optimum.

Submitted report, Sec. 4.1.3 and Table 4.3.

Heston adds a risk the stock cannot span

$$\begin{aligned}dS_t &= rS_t dt + \sqrt{v_t}S_t dW_t^S, \\dv_t &= \kappa(\theta - v_t) dt + \xi\sqrt{v_t} dW_t^v, \\dW_t^S dW_t^v &= \rho dt.\end{aligned}$$

$$dW_t^v = \rho dW_t^S + \sqrt{1 - \rho^2} dW_t^\perp$$

- ▶ v_t is instantaneous variance; $\sqrt{v_t}$ is volatility.
- ▶ Since $|\rho| < 1$, an orthogonal variance shock remains.
- ▶ With $\rho = -0.70$, the stock cannot span the full variance shock.

► Diffusion exposures

Instrument	dW^S	$dW^\perp$
stock	yes	no
liquid option	yes	yes

A liquid option with

$$\frac{\partial C^h}{\partial v} \neq 0$$

supplies a second locally independent exposure.

Scope: Local spanning holds in the ideal continuous-time frictionless model; discrete trading still leaves hedging error.

Source: Submitted report, Secs. 2.1.2 and 4.1.4; Heston (1993). Feller calculation in backup.

A liquid option supplies the missing variance exposure

$$G_T = \sum_{n=0}^{N-1} \delta_n \Delta S_n + \sum_{n=0}^{N-1} \eta_n \Delta C_n^h.$$

δ_n, η_n are total holdings; trades are $\Delta\delta_n, \Delta\eta_n$. Gains use holdings.

► Signed, bounded holdings

$$\delta_n = 5 \tanh(z_{1,n}), \quad \eta_n = 5 \tanh(z_{2,n}).$$

Holdings may be long or short; $(0, 1)$ is inappropriate. The bound stabilises training.

Source: Submitted report, Sec. 4.1.4 and Algorithm A.6.10.

► Why the hedge option differs

	Target	Hedge
Strike	$K = 0.9$	$K_h = 1.0$
Maturity	$T = 0.5$	$T_h = 1.0$
Role	liability	traded hedge

- $K_h = 1.0$ (ATM) gives strong variance sensitivity.
- $T_h > T$ keeps it alive throughout the hedge.

Each training sample is a full Heston trajectory



► Simulation step

$$Z_n^v = \rho Z_n^S + \sqrt{1 - \rho^2} Z_n^\perp.$$

$$S_{n+1} = S_n \exp \left[\left(r - \frac{1}{2} v_n^+ \right) \Delta t + \sqrt{v_n^+ \Delta t} Z_n^S \right].$$

$$\tilde{v}_{n+1} = v_n + \kappa(\theta - v_n^+) \Delta t + \xi \sqrt{v_n^+ \Delta t} Z_n^v,$$

$$v_{n+1} = \max(\tilde{v}_{n+1}, 0).$$

► Network state

$$x_n = \left(\log \frac{S_n}{K}, \frac{T - t_n}{T}, \frac{v_n}{\theta}, \log \frac{S_n}{K_h}, \frac{T_h - t_n}{T_h}, \frac{C_n^h}{S_n} \right).$$


- Log-Euler keeps $S_{n+1} > 0$.
- Full truncation stops negative variance from entering the next square root.
- One sample is a complete trajectory, not one isolated observation.

Boundary: Full-information simulation: true v_n enters the state; instantaneous variance is latent in real markets.

Source: Submitted report, Listing C.4 and Algorithms A.3.3, A.6.8, A.6.10.

A tempting shortcut: freeze the current volatility

$$C_t^h \approx C_{\text{BS}}(S_t, K_h, T_h - t, \sqrt{v_t}).$$



A plausible price level is not enough: the whole price path must be Heston-consistent.


Source: Submitted report, Sec. 4.1.4 (pp. 24–25) and Algorithm A.6.5.

- Why it seemed reasonable
 - cheap to evaluate;
 - responsive to current variance;
 - creates a variance-sensitive second instrument.
- Question Does repeatedly recomputing the Black–Scholes level produce a risk-neutral martingale under the simulated Heston dynamics?

The martingale test catches the inconsistency

Estimated mean total liquid-option movement over the target hedge horizon $[0, T]$

$$\hat{m} = \frac{1}{M} \sum_{i=1}^M \left(C_T^{h,(i)} - C_0^h \right), \quad T = 0.5 < T_h = 1.0.$$



► Proxy signal

$$\frac{-0.009482}{0.000459} \approx -20.6$$

standard errors below zero

- At $r = 0$, the martingale reference is zero mean movement.
- This does not mean each path is flat.
- The COS estimate is statistically negligible.

The proxy is not just imprecise in level; its path dynamics are inconsistent with the simulated Heston market.

Source: Submitted report, Table A.9 and Sec. 4.1.4. Diagnostic evidence, not a standalone arbitrage proof.

Why the frozen-volatility proxy drifts

Frozen-volatility liquid hedge-option proxy: $g(t, S, v) = C_{BS}(S, K_h, T_h - t, \sqrt{v})$

Apply Heston dynamics to the Black–Scholes proxy using Itô's lemma.

► Cancelled by the frozen- v Black–Scholes PDE

$$g_t + rSg_S + \frac{1}{2}vS^2g_{SS} - rg = 0$$

► Residual drift under Heston

$$R(t, S, v) = \kappa(\theta - v)g_v + \frac{1}{2}\xi^2vg_{vv} + \rho\xi vSg_{Sv}$$

► 1 variance mean reversion

► 2 variance curvature

► 3 spot–variance interaction

Repricing the Black–Scholes level at every date does not make the resulting option-price process satisfy the Heston pricing equation.

Scope: The residual is state-dependent; under the submitted simulation its average contribution was negative.

Source: Submitted report, Sec. 4.1.4 and App. A.6.3; Black–Scholes (1973); Heston (1993).

COS evaluates a Heston-consistent option price

$$C^H(t, S, v) = \mathbb{E}^{\mathbb{Q}}[e^{-r(T_h - t)}(S_{T_h} - K_h)^+ \mid S_t = S, v_t = v]$$

Heston supplies the characteristic function; COS maps each simulated Heston state to a liquid-option price and the sensitivities needed for hedging.

► Pricing and marking

Heston characteristic function →

COS price and sensitivities

► Exposure-matching positions

$$\eta_n = \frac{V_n^T}{V_n^h}, \quad \delta_n = \Delta_n^T - \eta_n \Delta_n^h$$

Monte Carlo Heston states → COS option marks →
hedge positions

► **Heston COS benchmark** positions from Heston-consistent sensitivities.


$V = \partial C / \partial v$. The option holding matches variance exposure; the stock holding removes the remaining stock exposure.

► **BS-proxy Greek heuristic** positions from frozen-volatility Black-Scholes sensitivities.

Same traded COS option path; different position rule.

Source: Submitted report, Sec. 4.1.4 and Apps. A.6.1–A.6.2; Heston (1993); Fang & Oosterlee (2008).

The learned option position adapts to the variance state



The stock-plus-option network broadly follows the local Heston COS delta-vega option position, while using less aggressive positions at higher variance.

Source: Submitted report, Figure 4.3 and Sec. 4.1.4. Vector crop from the final report; original export was not present in the submitted presentation repository.

The stock-plus-option NN beats the strongest tested analytic heuristic in all three seeds

Representative full-information run

Strategy	RMSE	Loss CVaR95
No hedge	0.190 061	0.492 959
BS-proxy stock delta	0.019 997	0.046 360
Heston COS delta-vega	0.014 782	0.033 251
Stock-only NN	0.018 166	0.041 379
BS-proxy Greek heuristic on the COS option path	0.009 874	0.023 130
Stock-plus-option NN	0.007 426	0.016 624

Heston COS delta-vega is the model-consistent local Greek benchmark; the BS-proxy rule is the strongest tested analytic heuristic.

Three independent reruns:
NN versus strongest analytic heuristic

Seed	NN	Heuristic	Imp. %
111	0.007 806	0.009 838	20.6
222	0.007 715	0.009 871	21.8
333	0.007 346	0.009 864	25.5

22.7%

mean RMSE improvement

27.1%

mean Loss CVaR95 improvement

Fair centring: Each strategy is fair-centred separately, so RMSE compares the remaining dispersion of hedge errors rather than differences in average error. See backup.

Interpretation: Model-consistent local Greek matching is not the same objective as finite-grid terminal-error minimisation.

Source: Submitted report, Tables 4.5–4.6 and A.11. Loss CVaR95 is mean seller loss in the worst 5% of test outcomes.

Three conclusions survive the diagnostics

► What the experiments establish

1. **Black–Scholes:** the neural strategy recovered the known finite-grid hedge.
2. **Conditioning:** adding the missing contract information repaired the identified generalisation failure.
3. **Heston:** with a consistently priced liquid option, the stock-plus-option NN achieved lower terminal error than the strongest tested analytic Greek heuristic in every submitted seed.


► What they do not establish

- high-powered statistical evidence beyond the three submitted seeds;
- a derived finite-grid optimal stock-and-option hedge;
- performance beyond the frictionless simulated Heston setting and submitted information assumptions.

A constrained neural policy can recover a known finite-grid solution and then improve terminal hedging performance when stochastic volatility adds a second risk factor and a second traded instrument.

Source: Submitted report, Secs. 4.1.4 and 5; stated limitations.


Backup: the learned Black–Scholes hedge matches the analytic curves



The learned hedge ratio closely follows the Black–Scholes delta and discrete-time MSE-optimal curves across moneyness.

Source: Submitted report, Figure 3.3 and Sec. 3.1.4. Vector crop from the final report.


Backup: training and validation losses converge without visible instability



Training and validation losses track closely and flatten late in training; the best validation loss occurs at the submitted epoch cap of 250.

Source: Submitted report, Figure 3.6 and Sec. 3.1.4. Optimisation and validation evidence, not proof of external generalisation.

State and tensor shapes



The same shared network is applied to every path-date state. Each complete path produces one terminal gain and one terminal hedge error.

Source: Submitted report, Listings C.1–C.4 and Algorithm A.6.10.

How gradients accumulate through the shared hedge

$$(\delta_{i,n}, \eta_{i,n}) = f_w(x_{i,n}), \quad L = B^{-1} \sum_{i=1}^B HE_i^2.$$
$$HE_i = \pi_w + \sum_{n=0}^{N-1} [\delta_{i,n} \Delta S_{i,n} + \eta_{i,n} \Delta C_{i,n}^h] - H_i.$$

$$\nabla_w L = \frac{2}{B} \sum_{i=1}^B HE_i \sum_{n=0}^{N-1} [\Delta S_{i,n} \nabla_w \delta_{i,n} + \Delta C_{i,n}^h \nabla_w \eta_{i,n}]$$

$$L \leftarrow HE \leftarrow G_T \leftarrow (\delta, \eta) \leftarrow f_w \leftarrow w$$

- ▶ one weight vector w is reused at every date;
- ▶ gradients accumulate across all paths and dates;
- ▶ no analytic delta or vega labels are supplied;
- ▶ COS prices are precomputed marks, outside the gradient path.

Source: Submitted report, Sec. 2.2.2 and Algorithm A.6.10.

Heston simulation

► Correlated shocks

$$Z_n^S, Z_n^\perp \sim N(0, 1), \quad Z_n^v = \rho Z_n^S + \sqrt{1 - \rho^2} Z_n^\perp.$$

► Log-Euler stock update

$$S_{n+1} = S_n \exp \left[\left(r - \frac{1}{2} v_n^+ \right) \Delta t + \sqrt{v_n^+ \Delta t} Z_n^S \right].$$

Preserves $S_{n+1} > 0$ and respects multiplicative stock dynamics.

► Truncation-based variance update

$$\begin{aligned} \tilde{v}_{n+1} &= v_n + \kappa(\theta - v_n^+) \Delta t + \xi \sqrt{v_n^+ \Delta t} Z_n^v, \\ v_{n+1} &= \max(\tilde{v}_{n+1}, 0), \quad v_n^+ = \max(v_n, 0). \end{aligned}$$

Ordinary Euler–Maruyama can propose negative variance, making the next square root invalid.

Boundary: Standard Monte Carlo generates many paths; log-Euler and full truncation specify each step. This is not an exact Heston transition. Feller positivity does not prevent negative Euler proposals. Discrete paths versus continuous-time COS marks contribute to the small residual COS drift.

Source: Heston (1993); Lord, Koekkoek & van Dijk (2010); submitted report, Alg. A.3.3 and Listing C.4.

Two-instrument gain construction

► Causal state and shared policy

$$x_n = \left(\log \frac{S_n}{K}, \frac{T - t_n}{T}, \frac{v_n}{\theta}, \log \frac{S_n}{K_h}, \frac{T_h - t_n}{T_h}, \frac{C_n^h}{S_n} \right)$$

$$(\delta_n, \eta_n) = 5 \tanh(f_w(x_n)).$$

► Terminal gain and training loss

$$G_T = \sum_{n=0}^{N-1} (\delta_n \Delta S_n + \eta_n \Delta C_n^h),$$

$$HE = \pi_w + G_T - (S_T - K)^+,$$

$$L = \mathbb{E}[HE^2].$$

Boundary: Executable notebooks and validation outputs were submitted separately; this slide maps the presentation equations to the submitted report algorithms and listings.

Source: Submitted report, Sec. 4.1.4 and Algorithm A.6.10.

Holdings, trades, and gains

► Network action: total holdings

$$(\delta_n, \eta_n) \quad \text{over } [t_n, t_{n+1}].$$

► Frictionless one-step gain

$$\delta_n \Delta S_n + \eta_n \Delta C_n^h.$$

► Trade increments

$$\Delta \delta_n = \delta_n - \delta_{n-1}, \quad \Delta \eta_n = \eta_n - \eta_{n-1}.$$

- the submitted gain equation uses holdings;
- trade increments determine transaction costs and turnover;
- calling δ_n or η_n a trade changes the implementation's meaning.

Source: Submitted report, Sec. 4.1.4 and Algorithm A.6.10.

Why the Heston outputs use scaled tanh

► Two signed outputs

$$\delta_n = 5 \tanh(z_{1,n}), \quad \eta_n = 5 \tanh(z_{2,n}).$$

$[-5, 5]$
numerical output range

► Why this map is used

- stock and option holdings can be long or short;
- sigmoid's $(0, 1)$ range is unsuitable;
- scaled tanh permits signed, bounded positions;
- the bound improves numerical stability, but is not an optimality theorem.

Submitted diagnostic: $\max |\eta^{NN}| \approx 1.33$, comfortably below five.

Source: Submitted report, Sec. 4.1.4, Algorithm A.6.10 and Table A.12.

Frozen-volatility drift derivation

For the liquid hedge-option proxy $g(t, S, v) = C_{BS}(S, K_h, T_h - t, \sqrt{v})$: $(dS)^2 = vS^2 dt$, $(dv)^2 = \xi^2 v dt$, $dS dv = \rho \xi v S dt$.

$$dg = \left[g_t + rSg_S + \kappa(\theta - v)g_v + \frac{1}{2}vS^2g_{SS} + \frac{1}{2}\xi^2vg_{vv} + \rho\xi vSg_{Sv} \right] dt$$

► **Frozen- v PDE cancellation** $+ g_S S \sqrt{v} dW^S + g_v \xi \sqrt{v} dW^v$.

$$g_t + rSg_S + \frac{1}{2}vS^2g_{SS} - rg = 0$$

► **Residual discounted drift**

$$\kappa(\theta - v)g_v + \frac{1}{2}\xi^2vg_{vv} + \rho\xi vSg_{Sv}$$

Boundary: At submitted $r = 0$, discounted and undiscounted residuals coincide. The sign is state-dependent; this diagnostic is not a standalone arbitrage proof.

Source: Submitted report, Sec. 4.1.4 and App. A.6.3; Black-Scholes (1973); Heston (1993).

Heston characteristic function used by COS

► Conditional transform of log price

$$\phi(u; \tau, x, v) = \mathbb{E}^{\mathbb{Q}}[e^{iuX_T} \mid X_t = x, v_t = v] = \exp(C(u, \tau) + D(u, \tau)v + iux), \quad X_T = \log S_T.$$

$$u_k = \frac{k\pi}{b-a}, \quad k = 0, \dots, N_{\text{COS}} - 1.$$

► 1

ϕ encodes the conditional distribution in transform space.

► 2

COS samples ϕ on a deterministic frequency grid.

► 3

The finite sum gives price, delta and $\partial C / \partial v$.

Source: Heston (1993); Fang & Oosterlee (2008); submitted report, Algorithm A.6.6.

COS pricing

Deterministic semi-analytical pricing: no nested Monte Carlo inside COS.

$$C_{\text{COS}} \approx e^{-r\tau} \sum_{k=0}^{N_{\text{COS}}-1} A_k V_k,$$

$$A_k = \frac{2}{b-a} \Re[\phi(u_k) e^{-i u_k a}], \quad u_k = \frac{k\pi}{b-a}.$$

The prime halves $k = 0$. A_k is a density coefficient; V_k is the analytic payoff coefficient. u_k is the k -th cosine frequency on $[a, b]$.

Heston supplies $\phi(u)$; COS evaluates the conditional expectation from $\phi(u_k)$.

- ▶ $N_{\text{COS}} = 256$, $[a, b] = [-7, 7]$;
- ▶ traded range $\tau_h \in [0.5, 1]$;
- ▶ Carr–Madan: $\alpha = 1.5$, $u \in [0, 200]$, 20,000 points;
- ▶ price discrepancy 5.1×10^{-8} ;
- ▶ finite-difference errors: delta 9.8×10^{-5} , variance sensitivity 4.7×10^{-4} .

Boundary: Validation is local to the submitted traded range.

Source: Heston (1993); Fang & Oosterlee (2008); Carr & Madan (1999); submitted report, Apps. A.6.1–A.6.2.

Analytic Heston stock-and-option hedge

$$\begin{aligned}dC^T &= \cdots + \Delta^T S \sqrt{v} dW^S + V^T \xi \sqrt{v} dW^v, \\d(\delta S + \eta C^h) &= \cdots + (\delta + \eta \Delta^h) S \sqrt{v} dW^S + \eta V^h \xi \sqrt{v} dW^v.\end{aligned}$$

Match $\eta V^h = V^T$ and $\delta + \eta \Delta^h = \Delta^T$:

$$\eta = \frac{V^T}{V^h}, \quad \delta = \Delta^T - \eta \Delta^h$$

► Heston COS delta–variance benchmark Heston COS supplies target and liquid hedge-option sensitivities.

► BS-proxy Greek heuristic on the COS option path Frozen- v BS sensitivities set positions; gains use the same COS marks.

Boundary: Local continuous-time first-order diffusion matching; not the finite-grid terminal-MSE optimum. The report table retains “delta–vega”, although $V = \partial C / \partial v$.

Source: Submitted report, Sec. 4.1.4 and Algorithms A.6.8–A.6.9.

Pricing accuracy and finite-grid hedge optimality are different questions

► What COS establishes

- model-consistent Heston option prices;
- independently validated prices and sensitivities;
- local first-order exposures for the COS Greek rule.

► What it does not optimise

- finite-grid terminal MSE;
- between-date gamma and variance curvature;
- spot–variance cross effects and changing Greeks.

Model misspecification may accidentally compensate for finite-grid effects in this tested parameter regime. That ordering does **not** imply that Black–Scholes is a better Heston pricer.

The neural policy trains directly on terminal MSE. Turnover is supporting diagnostic evidence, not proof of the ordering.

Source: Submitted report, Sec. 4.1.4 and Tables 4.5–4.6.

Fair-premium centring

$$\pi_j^* = \arg \min_{\pi} \mathbb{E}_{\text{test}} \left[(\pi + G_T^j - H)^2 \right] = \mathbb{E}_{\text{test}} [H - G_T^j]$$

$$HE_j = \pi_j^* + G_T^j - H, \quad \mathbb{E}_{\text{test}}[HE_j] = 0.$$

- ▶ applied separately to every headline strategy;
- ▶ removes test-sample mean error;
- ▶ RMSE compares centred dispersion; Loss CVaR95 compares centred downside tail risk;
- ▶ not used during training: the NN jointly learns a raw premium and hedge weights.

Source: Submitted report, Sec. 4.1.4, Tables 4.5–4.6 and A.11, and Algorithm A.6.10.

▶ Data roles remain separate

Training	Fit weights and learned premium
Validation	Select architecture / best epoch
Test	Final fair-centred comparison

Boundary: Because π_j^* is estimated on the same test paths as the metrics, this is not fully out-of-sample joint pricing and hedging.

Three-seed Heston comparison

► Comparator: BS-proxy Greek heuristic on the COS option path


Seed	NN RMSE	Heuristic RMSE	RMSE imp. (%)	CVaR95 imp. (%)
111	0.007 806	0.009 838	20.6	26.3
222	0.007 715	0.009 871	21.8	27.0
333	0.007 346	0.009 864	25.5	28.0
Mean pairwise improvement			22.7	27.1

- The three-seed cell retrained the stock-only tanh and stock-plus-option NNs and reevaluated the COS delta–variance and BS-proxy Greek analytic rules on each fresh test set.
- No hedge, BS-proxy stock delta and representative-only rows in the main table were not part of this four-strategy rerun.

Boundary: Three seeds demonstrate consistency across submitted runs; they do not support high-powered confidence intervals.

Source: Submitted report, Table A.11.

Backup: learned option position in the representative Heston run



► **Exact setting** Representative full-information COS run; horizontal axis is the Heston variance state (v), and the vertical axis is average liquid-option position.

► **Defensive interpretation** The NN broadly follows the local Greek hedge but softens at high variance, consistent with optimising terminal MSE rather than copying the analytic rule.

Source: Submitted report, Figure 4.3 and Sec. 4.1.4. In this run, $\max |\eta^{NN}| \approx 1.33$, inside the imposed $[-5, 5]$ range.

Observable-information robustness

► Causal variance proxy

- EWMA decay $\lambda = 0.94$;
- $\text{corr}(\hat{v}_t, v_t) = 0.777613$;
- relative variance-proxy RMSE = 0.369375.

► Submitted mean RMSE

Strategy	RMSE
Stock-plus-option NN	0.007 877
BS-proxy Greek heuristic on the COS option path	0.010 734

The EWMA state is a causal compressed history of past squared returns; the MLP itself remains non-recurrent.

Boundary: C_t^h / S_t remains contemporaneously observable and variance-informative, so this is not a stock-return-only filtering experiment.

Source: Submitted report, Sec. 4.1.4 and Tables A.10, A.12, A.15; J.P. Morgan/Reuters (1996).

Why the hedge option differs from the target

► Target liability

$$K = 0.9, \quad T = 0.5$$

► Liquid hedging option

$$K_h = 1, \quad T_h = 1$$

- near-ATM options generally have useful variance sensitivity;
- $T_h > T$ keeps the liquid hedging option alive throughout the target horizon;
- trading an identical target option would largely trivialise the problem;
- the target liability need not itself be a liquid traded contract.

Source: Submitted report, Sec. 4.1.4 and the target/hedge-option configuration.

What the martingale diagnostic measures

- ▶ Sample mean over $[0, T]$

$$\hat{m} = \frac{1}{M} \sum_{i=1}^M \left(C_T^{h,(i)} - C_0^h \right).$$

- ▶ Monte Carlo standard error

$$\widehat{\text{SE}}(\hat{m}) = \frac{s(C_T^h - C_0^h)}{\sqrt{M}}.$$

- ▶ individual paths fluctuate; the statistic concerns average movement;
- ▶ under $r = 0$, the reference is zero;
- ▶ the target horizon ends at $T = 0.5 < T_h = 1$;
- ▶ this diagnoses dynamic inconsistency, not standalone executable arbitrage.

Boundary: Proxy: -0.009482 (SE 0.000459); COS: -0.000118 (SE 0.000479). The COS residual is negligible, not exactly zero.

Source: Submitted report, Table A.9 and Algorithm A.6.8.

Backup: the discrete-time conditional projection

$$\min_{\pi, \delta} \mathbb{E} \left[\left(\pi + \sum_{j=0}^{N-1} \delta_j \Delta S_j - H \right)^2 \right],$$
$$0 = \mathbb{E} \left[\left(H - \pi - \sum_{j=0}^{N-1} \delta_j \Delta S_j \right) \Delta S_n \mid \mathcal{F}_{t_n} \right].$$

Under submitted $r = 0$, S is a martingale, so

$$\delta_n^{\text{DT}} = \frac{\mathbb{E}[H \Delta S_n \mid \mathcal{F}_{t_n}]}{\mathbb{E}[(\Delta S_n)^2 \mid \mathcal{F}_{t_n}]} = \frac{\text{Cov}(H, \Delta S_n \mid S_{t_n})}{\text{Var}(\Delta S_n \mid S_{t_n})}.$$

Boundary: This simplification is tied to the submitted $r = 0$ martingale setting; local and global quadratic hedges need not coincide more generally.

Source: Submitted report, Appendix A.4; Schweizer (1995, 2001).

Backup: method-to-code map

Mathematical concept	Displayed implementation	Report evidence	Primary expert
Exact BS path simulation	<code>simulate_gbm_paths</code>	Listing C.1; Alg. A.3.1	Micaela / Glasson
Self-financing terminal error	<code>hedge_error</code>	Listing C.2; Alg. A.6.1	Micaela / Glasson
Shared Markov hedge	<code>SharedMarkovHedge</code>	Listing C.3; Sec. 2.2.2	Kaiden / Micaela
Heston full-truncation simulation	PyTorch loop in Listing C.4	Listing C.4; Alg. A.3.3	Glasson / Siphelele
COS liquid-option pricing	Algorithms A.6.6–A.6.8	App. A.6.1–A.6.2	Glasson
Two-instrument neural gain	Algorithm A.6.10	Sec. 4.1.4	Siphelele / Glasson

Boundary: Executable notebooks and validation outputs were submitted separately; this slide maps the presentation equations to the submitted report algorithms and listings.

Backup: transaction costs change the learned policy

$$V_T^\lambda = \pi + \sum_{n=0}^{N-1} \delta_n \Delta S_n - \lambda \sum_{n=0}^{N-1} S_{t_n} |\delta_n - \delta_{n-1}|, \quad \delta_{-1} = 0.$$

λ	Strategy	RMSE	Loss CVaR95	Turnover
0.0025	BS cost-adjusted	0.008 492	0.021 464	4.497 364
0.0025	Cost-aware NN	0.008 077	0.018 694	4.386 880
0.0050	BS cost-adjusted	0.010 013	0.025 565	4.497 364
0.0050	Cost-aware NN	0.008 331	0.019 269	4.284 602

Boundary: Comparator is a cost-adjusted frictionless delta, not a full optimal-cost benchmark such as Leland or Whalley–Wilmott.

Source: Submitted report, Appendix A.7 and Table A.16.

Backup: data-generation comparison

Paths	Method	RMSE	Loss CVaR95
10,000	Crude MC	0.009 430	0.021 266
10,000	LHS	0.009 201	0.020 845
10,000	Sobol + Brownian bridge	0.009 174	0.020 865
100,000	Crude MC	0.007 865	0.018 218
100,000	LHS	0.007 846	0.018 331
100,000	Sobol + Brownian bridge	0.007 840	0.018 125

At 10,000 paths, structured methods reduce RMSE by roughly 2.4–2.7%; by 100,000 paths the difference is negligible.

Source: Submitted report, Sec. 3.1.5 and Tables A.7–A.8.

Backup: experimental protocol

Experiment	Train	Val	Test	Batch	Training / seeds
Architecture selection	100k	20k	50k	256	60 epochs; 2026–2028
Final BS benchmark	100k	30k	100k	4096	250 epochs; base 123
Parameter robustness	50k	15k	50k	4096	150 epochs
Conditioned hedger	200k	30k	50k / case	4096	250 epochs; 777
Heston stock + option	100k	25k	100k	4096	150 epochs; base 123

- ▶ Adam learning rate 10^{-3} unless otherwise stated.
- ▶ Baseline: $S_0 = 1, K = 0.9, T = 0.5, r = 0, \sigma = 0.4, N = 125$.
- ▶ Architecture search uses more updates per epoch via batch size 256; final experiments use 4096.

Source: Submitted report, Appendix A.1, Table A.1.

Backup: AI assistance and verification

Assistance area disclosed	Human verification reported
Code drafting and debugging	Executed notebooks, code inspection and exported results
Mathematical explanations and LaTeX	Checked against theory and cited sources
Report organisation and consistency	Tables checked against CSV/TeX outputs and figures
COS implementation support	Independent Carr–Madan Fourier price check
Neural experiment support	BS analytic and finite-grid benchmarks; archived outputs
Documentation and reproducibility	Checklists, packaging review and author decisions

Boundary: The report names ChatGPT and Claude. It states that AI suggestions were not authoritative and that all final methodological and interpretive decisions remained with the authors.

Source: Submitted report, Appendix B.

Backup: architecture selection by seed

Seed	Validation winner	Val. loss	Test RMSE
2026	Shared norm. 64×3 , tanh	6.025×10^{-5}	0.007 883
2027	Shared norm. 64×3 , tanh	6.164×10^{-5}	0.007 916
2028	Shared norm. 64×2 , ReLU	6.014×10^{-5}	0.007 855

The 64×3 tanh model won two seeds and had the best mean validation loss and test RMSE; the ReLU alternative was close.

Source: Submitted report, Tables A.3–A.6.

Backup: references I

- Fischer Black and Myron Scholes. The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3):637–654, 1973.
- Hans Buehler, Lukas Gonon, Josef Teichmann, and Ben Wood. Deep hedging. *Quantitative Finance*, 19(8):1271–1291, 2019.
- Peter Carr and Dilip B. Madan. Option valuation using the fast fourier transform. *Journal of Computational Finance*, 2(4):61–73, 1999.
- Fang Fang and Cornelis W. Oosterlee. A novel pricing method for european options based on fourier-cosine series expansions. *SIAM Journal on Scientific Computing*, 31(2):826–848, 2008.
- Hans Föllmer and Martin Schweizer. Hedging of contingent claims under incomplete information. In M. H. A. Davis and R. J. Elliott, editors, *Applied Stochastic Analysis*, pages 389–414. Gordon and Breach, 1991.
- Jim Gatheral. *The Volatility Surface: A Practitioner's Guide*. Wiley, 2006.
- Paul Glasserman. *Monte Carlo Methods in Financial Engineering*. Springer, 2004.
- Steven L. Heston. A closed-form solution for options with stochastic volatility. *The Review of Financial Studies*, 6(2):327–343, 1993.
- J.P. Morgan/Reuters. Riskmetrics technical document. Technical report, J.P. Morgan/Reuters, 1996.
- Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *International Conference on Learning Representations*, 2015.
- Hayne E. Leland. Option pricing and replication with transactions costs. *Journal of Finance*, 40(5):1283–1301, 1985.
- Roger Lord, Remmert Koekkoek, and Dick van Dijk. A comparison of biased simulation schemes for stochastic volatility models. *Quantitative Finance*, 10(2):177–194, 2010.
- Robert C. Merton. Theory of rational option pricing. *The Bell Journal of Economics and Management Science*, 4(1):141–183, 1973.
- Johannes Ruf and Weiguan Wang. Neural networks for option pricing and hedging: A literature review. *Journal of Computational Finance*, 24(1):1–46, 2020.
- Martin Schweizer. Variance-optimal hedging in discrete time. *Mathematics of Operations Research*, 20(1):1–32, 1995.
- Martin Schweizer. A guided tour through quadratic hedging approaches. In Elyes Jouini, Jakša Cvitanić, and Marek Musiela, editors, *Option Pricing, Interest Rates and Risk Management*, pages 538–574. Cambridge University Press, 2001.
- A. E. Whalley and P. Wilmott. An asymptotic analysis of an optimal hedging model for option pricing with transaction costs. *Mathematical Finance*, 7(3):307–324, 1997.